



Test-Set Confusion Matrices Across Core Architectures (N = 320)


(a) Linear SVM
[Traditional ML] Macro-F1: 0.9917




(b) Multinomial Naive Bayes
[Traditional ML] Macro-F1: 1.0000




(c) Random Forest
[Traditional ML] Macro-F1: 1.0000




(d) CNN-BiLSTM
[Deep Learning] Macro-F1: 0.9959



(e) Attention-BiLSTM
[Deep Learning] Macro-F1: 0.9836



(f) mBERT
[Transformer] Macro-F1: 0.9836



Non-Offensive (0) Offensive (1)

Predicted Label

Non-Offensive (0) Offensive (1)

Predicted Label

Non-Offensive (0) Offensive (1)

Predicted Label